

Student Score - Supervised ML - The Spark Foundation

July 15, 2022

1 Student Score Prediction Using Supervised ML

```
[29]: # Importing all libraries required in this notebook
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
[30]: import warnings
warnings.filterwarnings('ignore')
```

```
[31]: #Uploading dataset

data = "https://raw.githubusercontent.com/AdiPersonalWorks/Random/master/
↪student_scores%20-%20student_scores.csv"
student_data = pd.read_csv(data)
student_data.head(10)
```

```
[31]:
```

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30
5	1.5	20
6	9.2	88
7	5.5	60
8	8.3	81
9	2.7	25

```
[32]: student_data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25 entries, 0 to 24
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -

```

```
0    Hours    25 non-null    float64
1    Scores  25 non-null    int64
dtypes: float64(1), int64(1)
memory usage: 528.0 bytes
```

```
[33]: student_data.describe()
```


```
[33]:
```

	Hours	Scores
count	25.000000	25.000000
mean	5.012000	51.480000
std	2.525094	25.286887
min	1.100000	17.000000
25%	2.700000	30.000000
50%	4.800000	47.000000
75%	7.400000	75.000000
max	9.200000	95.000000

```
[34]: #Investigating the correlation between the columns


sns.pairplot(student_data)
```

```
[34]: <seaborn.axisgrid.PairGrid at 0x1f89cf95cd0>
```



```
[35]: # Scatter plot


student_data.plot(x='Hours', y='Scores' , style='o')
plt.title('Hours vs Score percentage')
plt.xlabel('Hours Studied')
plt.ylabel('Percentage Score')
plt.show()
```



```
[36]: #Checking the column to be predicted on how its is distributed
```

```
sns.distplot(student_data['Hours'])
```


```
[36]: <AxesSubplot:xlabel='Hours', ylabel='Density'>
```



From the above plot we can see the average distribution is around 1 to 3 hours and the second highest is between 7 to 9

```
[37]: #check the correlation between each columns using a heat map  
  
sns.heatmap(student_data.corr(), annot = True )
```

```
[37]: <AxesSubplot:>
```



2 Assigning Test and Training sets

```
[38]: #lets assign X label

X = student_data.iloc[:, :-1].values

#Lets assign the target variable 'Hours' for the Y label

y = student_data.iloc[:, 1].values
```

```
[39]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=0)
```

3 Training the model

```
[40]: from sklearn.linear_model import LinearRegression

LM = LinearRegression()

LM.fit(X_train, y_train)
```

```
[40]: LinearRegression()
```

4 Making Predictions

```
[41]: print(X_test) # Testing data - In Hours
```

```
[[1.5]
 [3.2]
 [7.4]
 [2.5]
 [5.9]]
```

```
[42]: Pred = LM.predict(X_test) # giving the model an unexpected data to predict
```

```
[43]: Pred #this shows the prediction made by the model for the X_test data
```

```
[43]: array([16.88414476, 33.73226078, 75.357018 , 26.79480124, 60.49103328])
```

```
[44]: # Comparing Actual vs Predicted
```

```
new_df = pd.DataFrame({'Actual': y_test, 'Predicted': Pred})
new_df
```

```
[44]:
```

	Actual	Predicted
0	20	16.884145
1	27	33.732261
2	69	75.357018
3	30	26.794801
4	62	60.491033

```
[45]: #now lets check the intercepts by calling the model
```

```
print(LM.intercept_)
```

```
2.018160041434683
```

```
[46]: #the o/p coefficient will relate to each and every value in the dataset
```

```
LM.coef_
```


```
[46]: array([9.91065648])
```

```
[47]: # regression line
```

```
Reg_line = LM.coef_*X+LM.intercept_
```

```
#Lets check the above in visualization
```


```
plt.scatter(X, y)
plt.plot(X, Reg_line);
plt.show()
```



We can see that both the `y_test` data and the `X_train` data are more or less near the linear line

```
[48]: #Let plot a histogram of the residuals if any  
  
sns.distplot((y_test - Pred)) #Residual is the difference between the actual_  
    ↪value and the predicted value
```

```
[48]: <AxesSubplot:ylabel='Density'>
```

5 Evaluation of Regression

Now our goal is to minimise the error if any with various evaluation metrics

```
[49]: from sklearn import metrics
```

```
[50]: #Mean Average Error (MAE) - its the average of errors
print('Mean Average Error is :', metrics.mean_absolute_error(y_test,Pred))
```

Mean Average Error is : 4.183859899002975

```
[51]: #Mean Squared Error (MSE) - better cancellation of larger error compared to MAE
      ↪as it is squaring them
print('Mean Squared Error is :', metrics.mean_squared_error(y_test,Pred))
```

Mean Squared Error is : 21.5987693072174

```
[52]: #Root Mean Squared Error (RMSE) - Most accurate of all as it can be interpreted
      ↪in the 'y' units
print('Root Mean Squared Error is :', np.sqrt(metrics.
      ↪mean_squared_error(y_test,Pred)))
```

Root Mean Squared Error is : 4.6474476121003665

We can consider RMSE evaluation metric for our model as its the most accurate out of the 3

6 Now Lets Test the Model with our own data to Predict the Score

```
[53]: input_value = np.array(9.25).reshape(-1, 1)
      Pred = LM.predict(input_value)
      print("No of Hours = {}".format(input_value))
      print("Predicted Score = {}".format(Pred[0]))
```

```
No of Hours = [[9.25]]
Predicted Score = 93.69173248737538
```

```
[ ]:
```